

JustInTimeout: Enacting Dynamic Network Timeouts via Distributed Feedback

Yash Trivedi
Dylan Wilson
Nisarg Chokshi

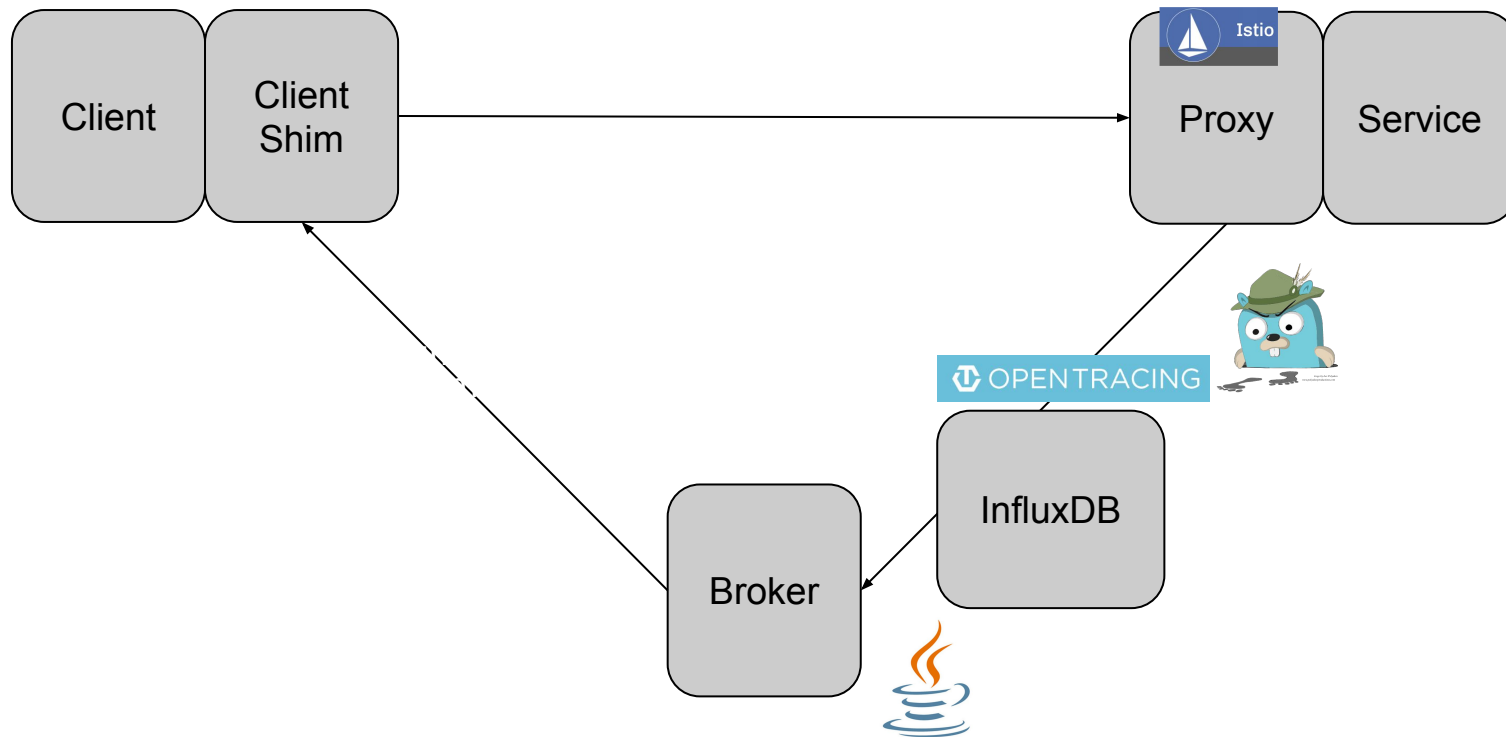
Agenda

- Big picture
 - Goals
 - Logical View
- System Overview
 - RNN Broker
 - Client Shim
 - GP Broker
- Results
- Future Directions

Goals

- Use network response data to predict near-term timeout values.
- Transparently inject appropriate timeout values into network calls.

JustInTimeout: Logical View



Recurrent Neural Network Broker

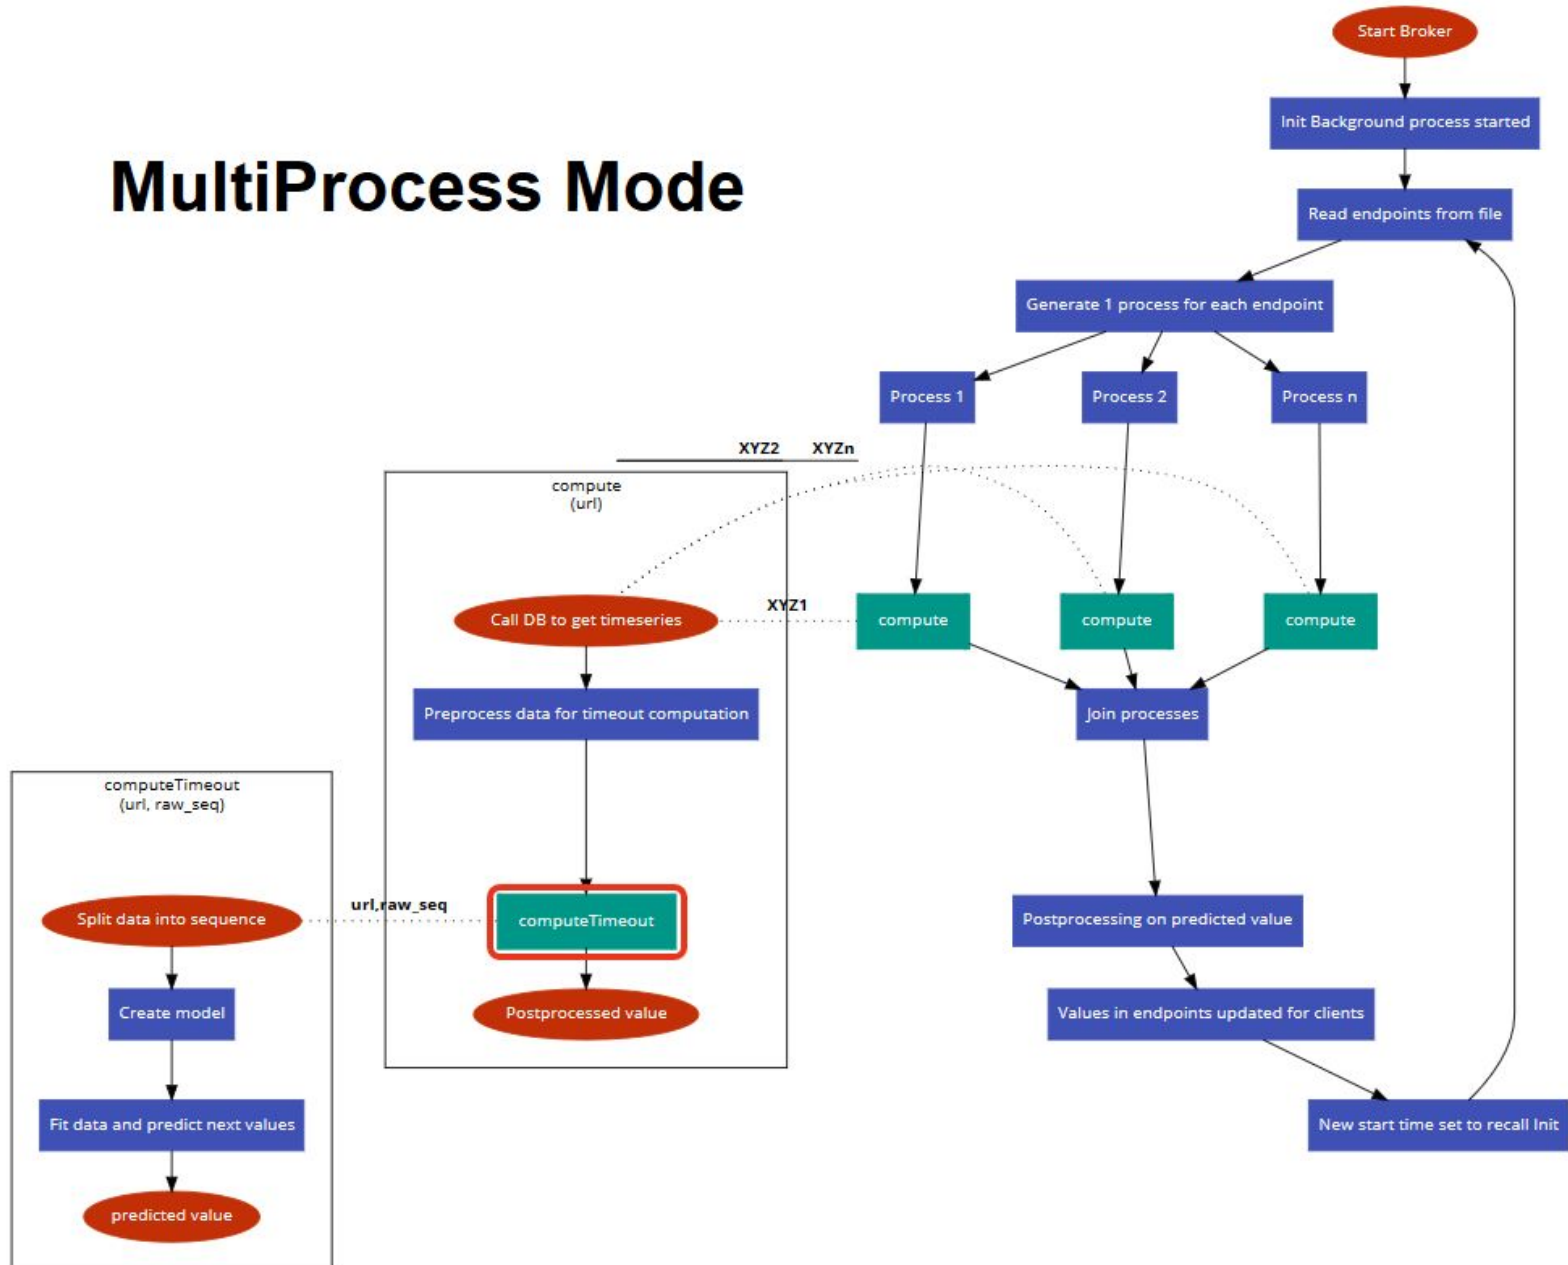
Why Recurrent Neural Network

- This problem can be simplified as a **time series prediction problem** where the previous inputs within a short window span will most likely be the best candidates to predict future value.
- **Convolved LSTM** is optimal when you have no need for convolution operation in the input data all the while making use of CNNs along with LSTM for fast predictions.

MultiProcess vs MultiDocker

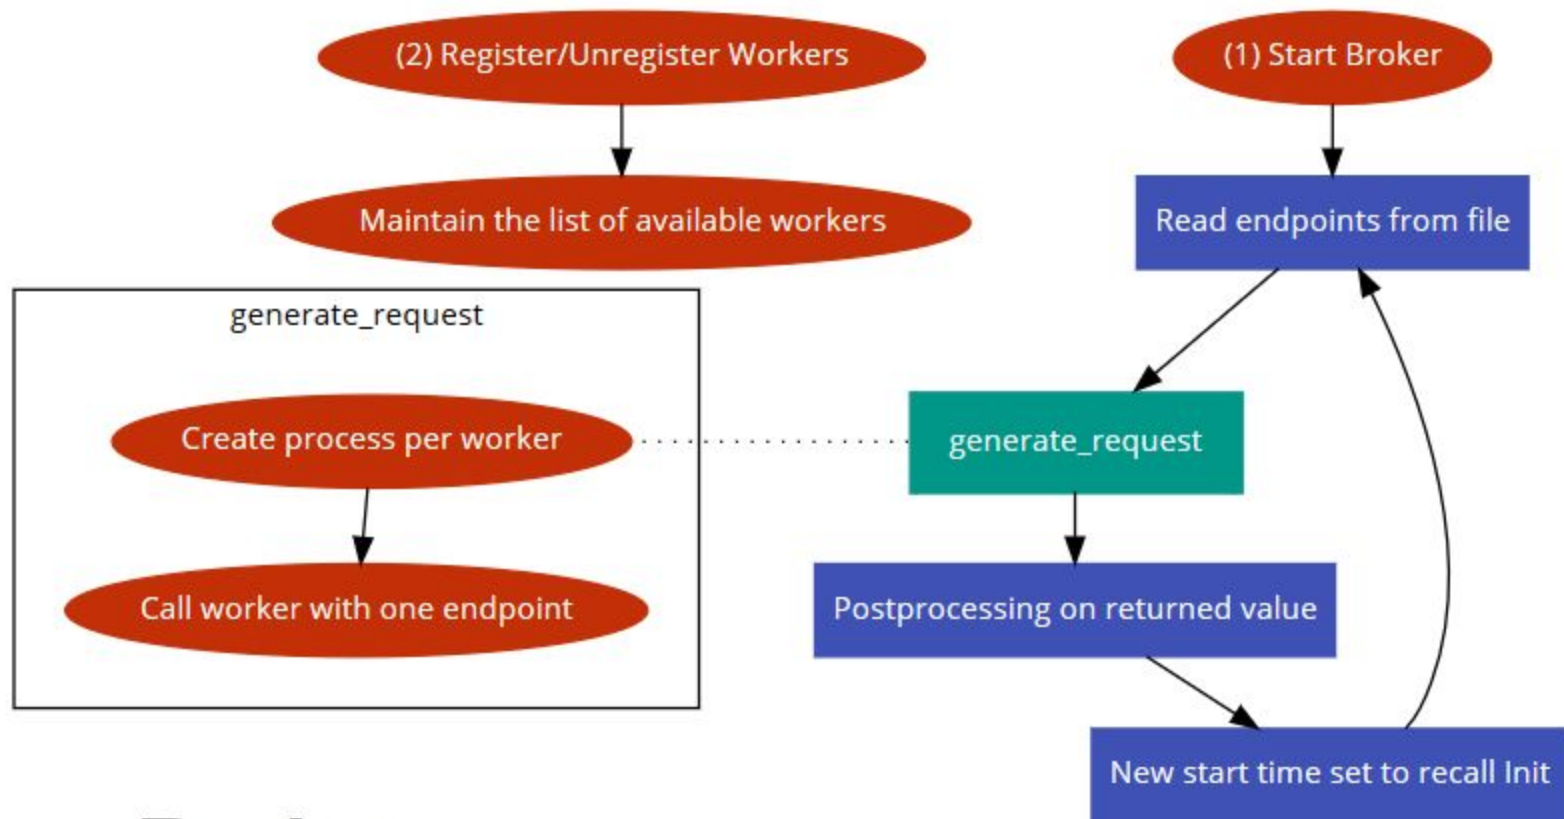
- In MultiProcess mode, there is just one server running which spawns off multiple processes which individually do the timeout computation for the url assigned to them.
- In the MultiDocker mode, this idea is taken to the extreme where there are workers pods who register themselves with the running broker and get tasks.

MultiProcess Mode

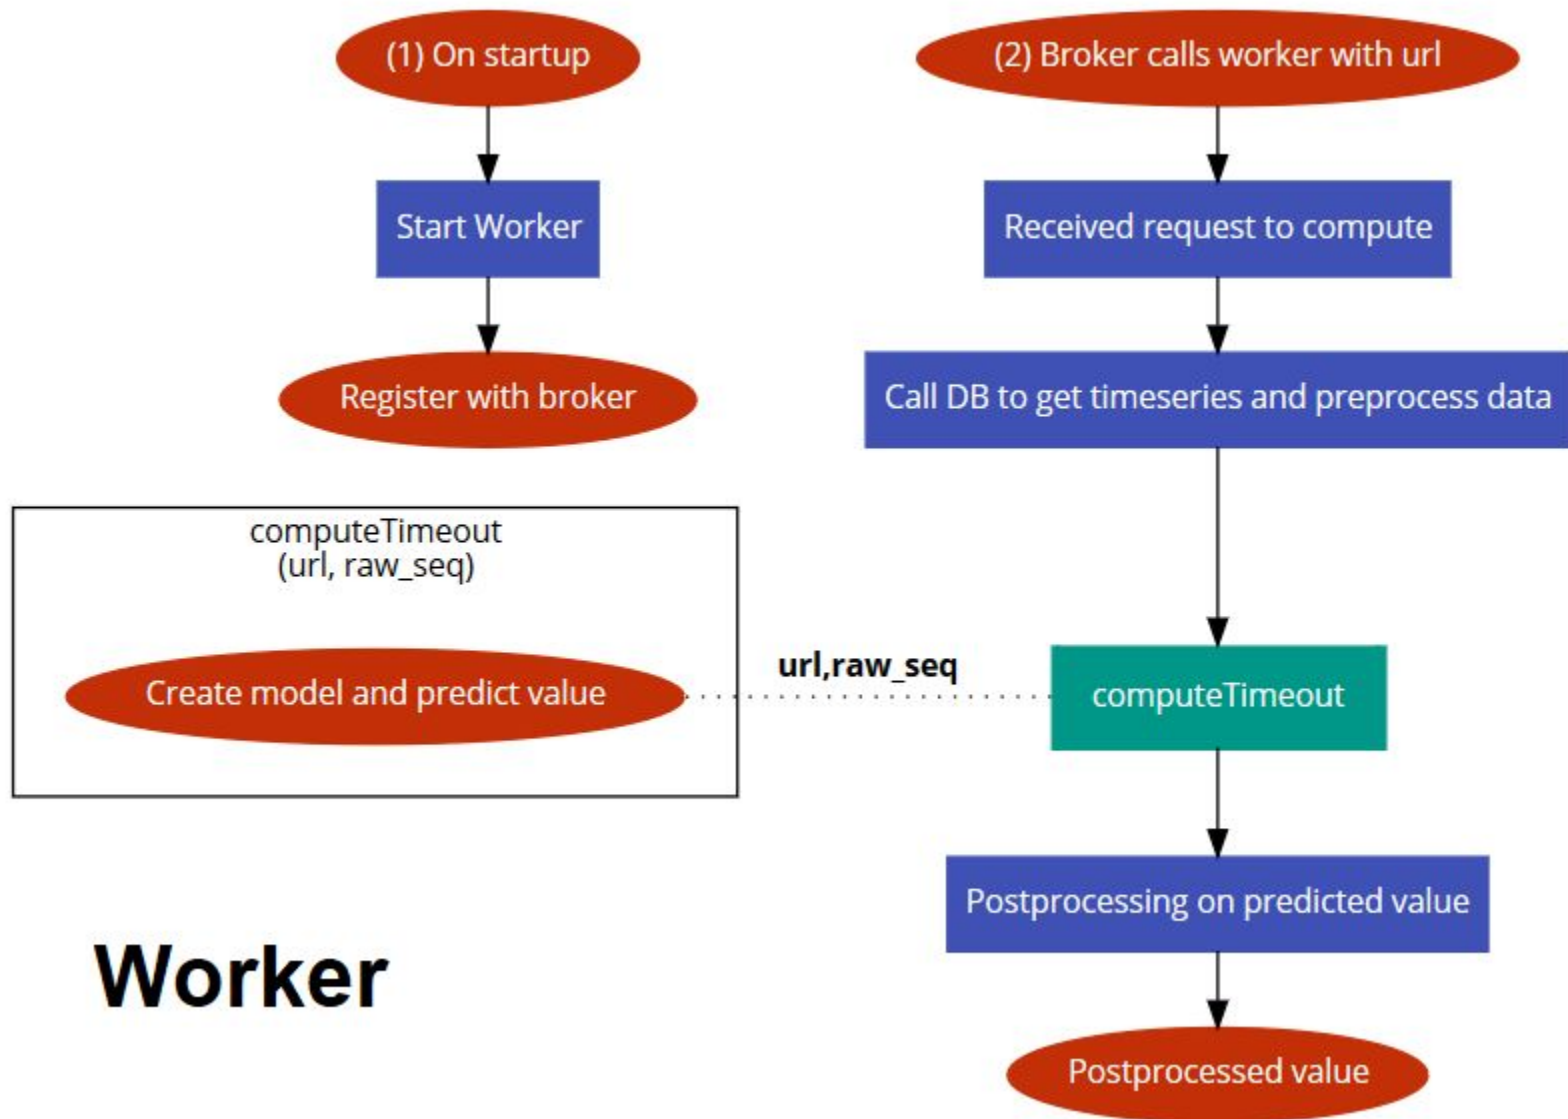


Enhancements in MultiDocker

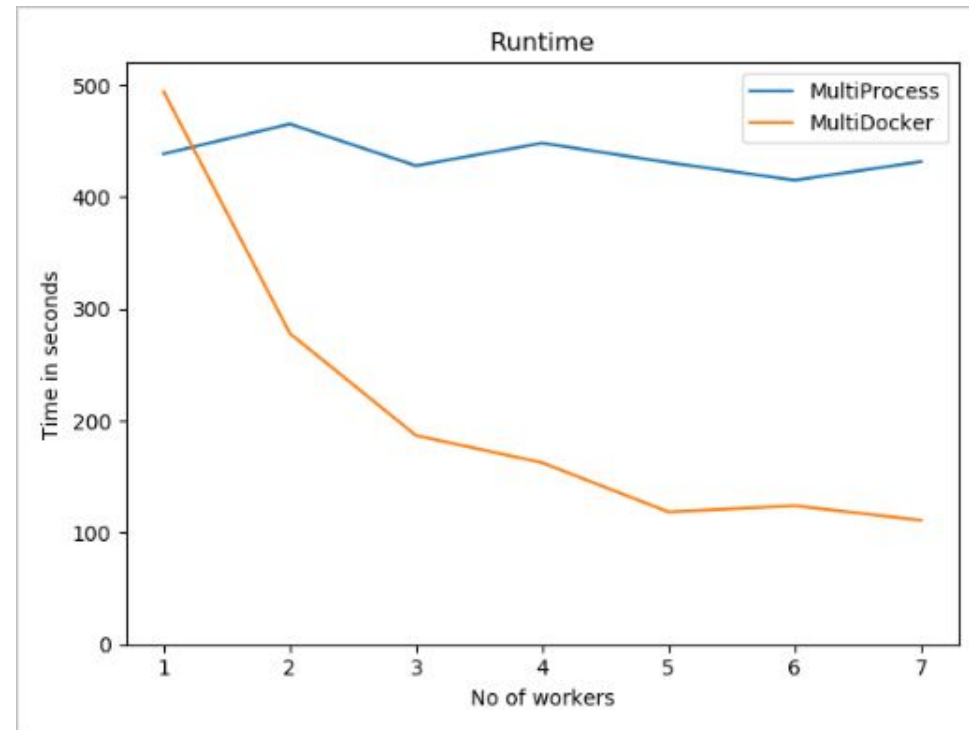

- This approach requires workers to register with the broker.
- Two new url calls which workers can use to register and unregister from the broker.
- Once registered, they will carry out the computations.



Broker



Performance

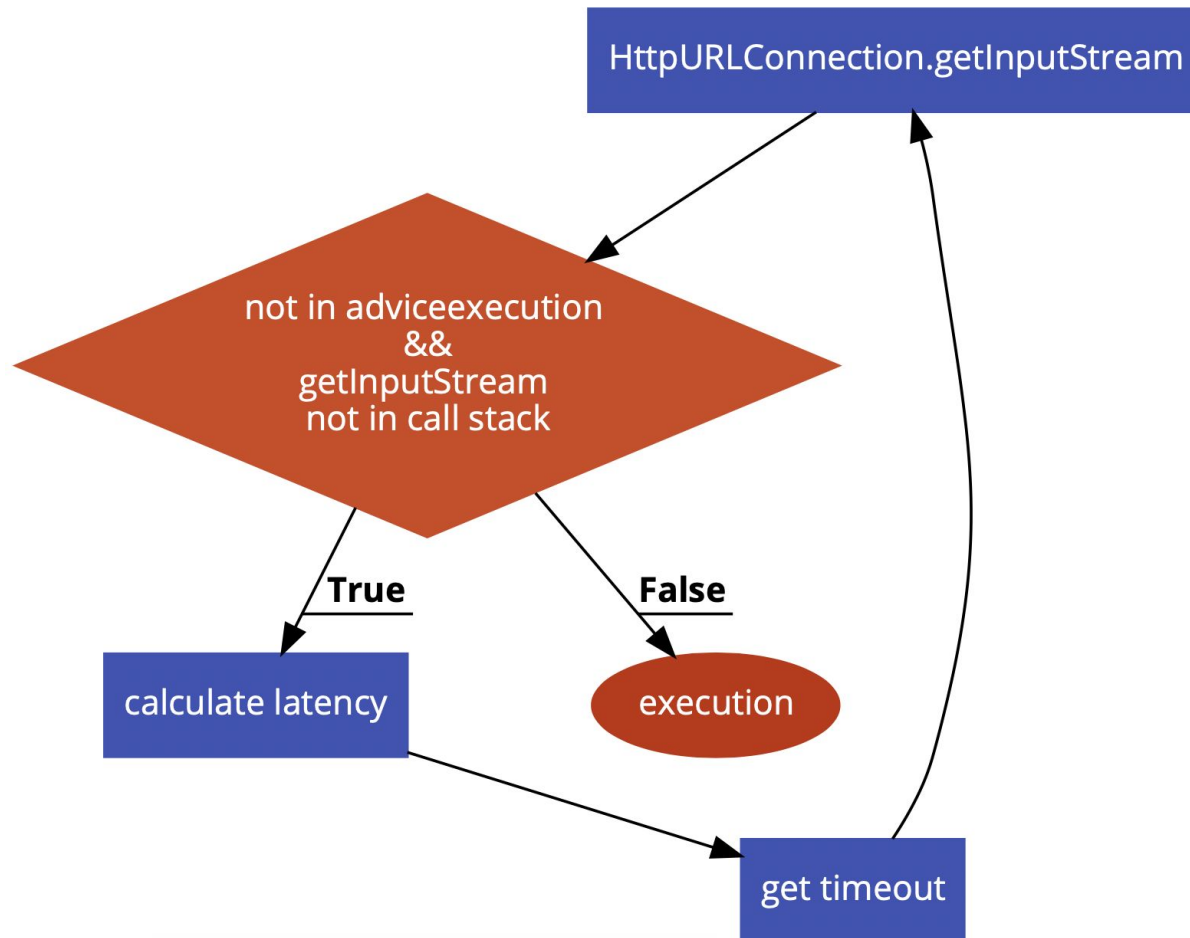


Client Shim

Aspect Oriented Timeout Injection

- Distributed feedback for timeout adjustment
 - Fetching Timeouts
 - Caching Timeouts
 - Calculating Latency
- Weaved into the JRE



















Advicing the Advice



Portability of the Shim

- Boot class path
- System properties
- Environment Variables
- Dynamically disable injection

Timeout Injection

```
▶  2020-04-28 20:22:45.294 CDT java.net.SocketTimeoutException: Read timed out at sun.reflect.NativeConstructorAccessorImpl.  
▶  2020-04-28 20:23:07.668 CDT java.net.SocketTimeoutException: Read timed out at java.net.SocketInputStream.socketRead0(Nat  
▶  2020-04-28 20:23:07.668 CDT There was an error processing request number 2  
▶  2020-04-28 20:23:07.668 CDT java.net.SocketTimeoutException: Read timed out at sun.reflect.NativeConstructorAccessorImpl.  
▶  2020-04-28 20:24:36.538 CDT java.net.SocketTimeoutException: Read timed out at java.net.SocketInputStream.socketRead0(Nat  
▶  2020-04-28 20:24:36.539 CDT There was an error processing request number 10  
▶  2020-04-28 20:24:36.539 CDT java.net.SocketTimeoutException: Read timed out at sun.reflect.NativeConstructorAccessorImpl.  
▶  2020-04-28 20:25:55.664 CDT java.net.SocketTimeoutException: Read timed out at java.net.SocketInputStream.socketRead0(Nat  
▶  2020-04-28 20:25:55.665 CDT There was an error processing request number 17  
▶  2020-04-28 20:25:55.665 CDT java.net.SocketTimeoutException: Read timed out at sun.reflect.NativeConstructorAccessorImpl.  
▶  2020-04-28 20:26:05.354 CDT java.net.SocketTimeoutException: Read timed out at java.net.SocketInputStream.socketRead0(Nat  
▶  2020-04-28 20:26:05.355 CDT There was an error processing request number 18  
▶  2020-04-28 20:26:05.355 CDT java.net.SocketTimeoutException: Read timed out at sun.reflect.NativeConstructorAccessorImpl.  
▶  2020-04-28 20:26:49.899 CDT java.net.SocketTimeoutException: Read timed out at java.net.SocketInputStream.socketRead0(Nat  
▶  2020-04-28 20:26:49.902 CDT There was an error processing request number 22  
▶  2020-04-28 20:26:49.902 CDT java.net.SocketTimeoutException: Read timed out at sun.reflect.NativeConstructorAccessorImpl.  
▶  2020-04-28 20:26:58.813 CDT java.net.SocketTimeoutException: Read timed out at java.net.SocketInputStream.socketRead0(Nat  
▶  2020-04-28 20:26:58.816 CDT There was an error processing request number 23
```


Gaussian Process Broker

Bayesian Synthesis of Gaussian Processes

Sample from the set of possible GP programs to find a set the likely generated the data.


Iteratively move from a vaguely random set of programs to a likely set of programs.

Gaussian processes




Program Synthesis

Time series structure discovery using program synthesis in Venture



```
(+ (* (WHITE-NOISE 49.5)
      (CONSTANT 250.9))
  (+ (PERIODIC 13.2 8.6)
      (+ (LINEAR 1.2)
          (LINEAR 4.9)
          (WHITE-NOISE 0.1)))
```



```
// ** LOAD THE TIME SERIES DATA **
```

```
define xs_obs = get_data_xs("./data-train.csv");
```

```
define ys_obs = get_data_ys("./data-train.csv");
```

```
define xs_test = get_data_xs("./data-test.csv");
```

```
define ys_test = get_data_ys("./data-test.csv");
```

```
// ** SAMPLE ENSEMBLE OF GAUSSIAN PROCESSES FROM THE PRIOR **
```

```
resample(100);
```

```
assume dsl_code ~ generate_random_program();
```

```
assume gaussian_process_model = produce_gaussian_process(dsl_code);
```

```
observe gaussian_process_model({xs_obs}) = ys_obs;
```

```
// ** RUN BAYESIAN SYNTHESIS **
```

```
repeat(1000, {
```

```
  infer resimulation_mh(/?hypers/);
```

```
  infer resimulation_mh(/?structure/))}
```

```
// ** OBTAIN FORECASTS **
```

```
sample_all(gaussian_process_model({xs_test}))
```

GP Model Exploration

Scenarios:

1. Long Ramp
2. High Low Cycles
3. Random

Various Ensemble Counts

Various Iterations

~2000 seconds of training data

~1200 seconds forward prediction

```
CREATE CONTINUOUS QUERY "cq_20s_max" ON "tracing"  
BEGIN  
  SELECT max("duration") INTO "max_duration" FROM "span"  
  GROUP BY time(20s), service_name  
END
```

Model Exploration: Random

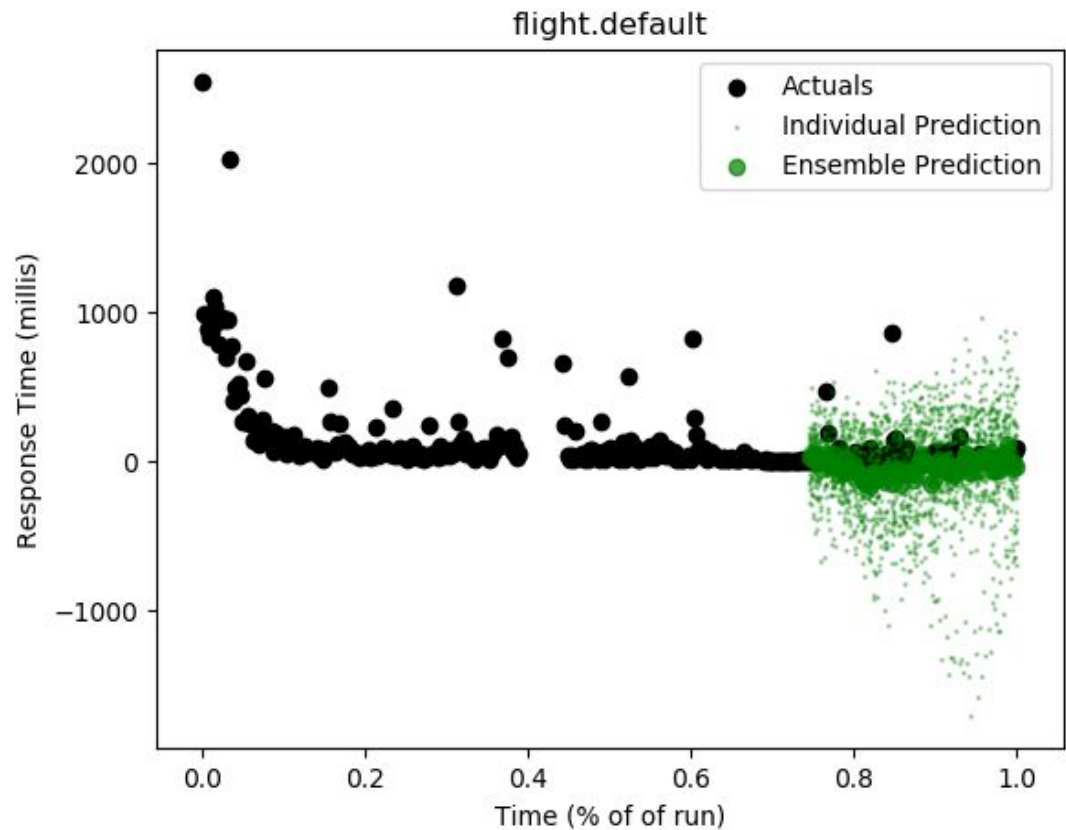
Ensemble: 100

Iterations: 1000

Train Time: 5544 s

Train Data: 90 m

Under Predict: 81%



Model Exploration: Cyclical

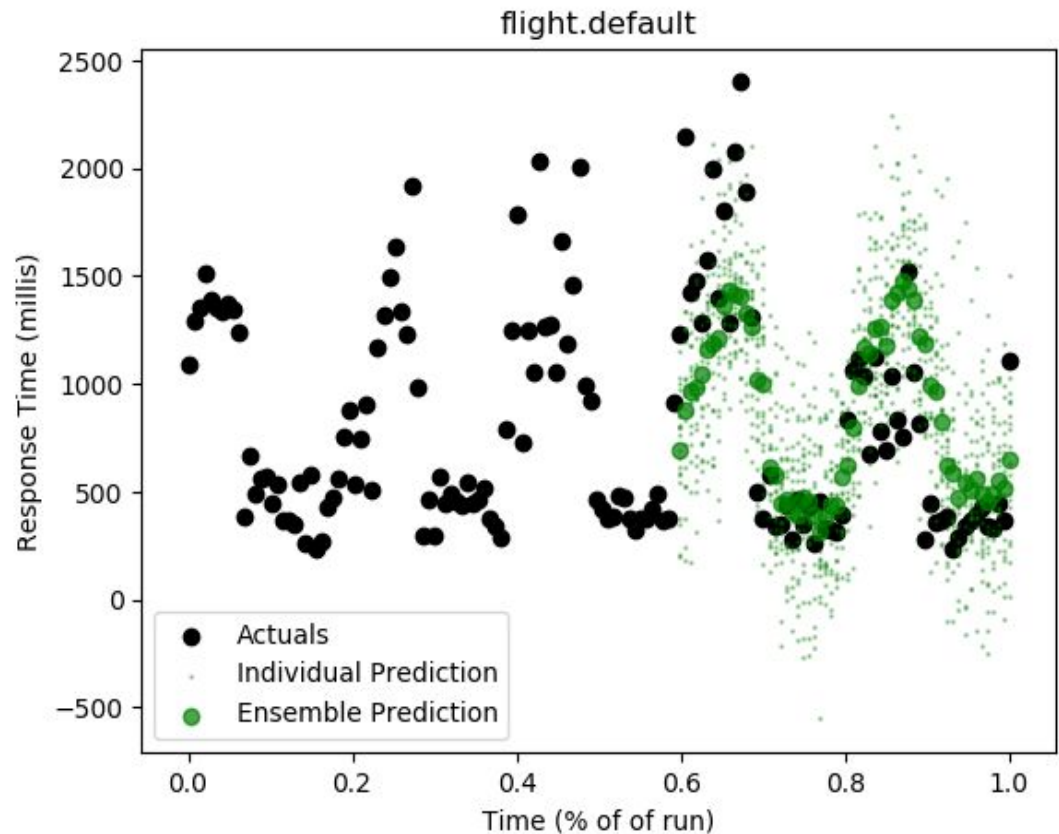
Ensemble: 20

Iterations: 1000

Train Time: 576 s

Train Data: 31 m

Under Predict: 22%



Model Exploration: Ramp

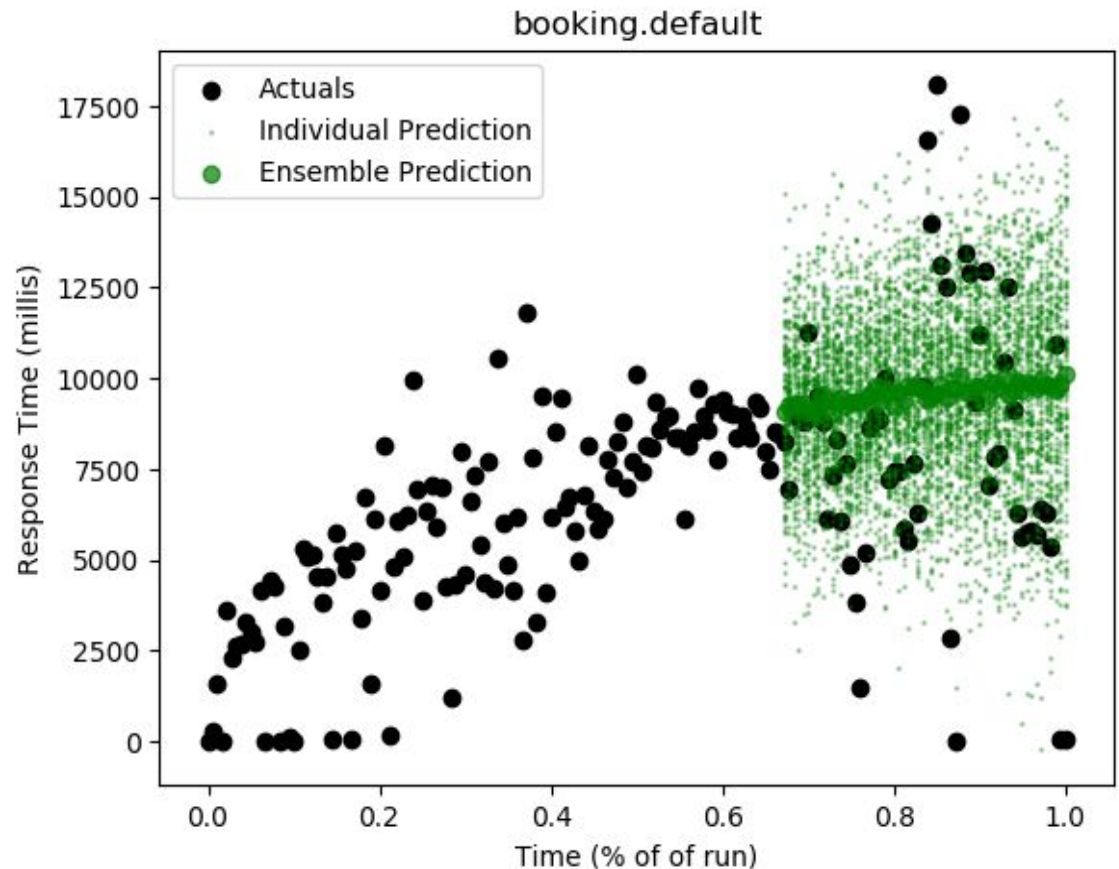
Ensemble: 100

Iterations: 100

Train Time: 359 s

Train Data: 40 m

Under Predict: 28%



GP Setup For Broker Use

- 20 second granularity using max
- (Train) on \leq last 10 minutes - Ensemble Count = 1
- Predict ~7 minutes
- Timeout Recommended = 150% of
 - 99th Percentile of entire predicted next 7 minutes

GP Future Directions

- Unify the Broker version with the experiments. Run parameter optimization experiments.
- Predict over multiple time granularities and combine
- Use uncertainty measures to drive multipliers of timeout values
- Run syntactic analysis of the generated programs to find change points and trigger alarms or auto-scale resources

Integrated System Evaluation

Repeating cycles of lower and higher load over ~ 90 minutes.

4 GKE Nodes

3 Prob Broker

1 RNN Broker

Client Data Collector - “YClient” sampling the endpoints and the brokers.

Results

Table 1: Performance Metrics of Each Technique by Endpoint

Endpoint	RNN		GP	
	Under-predict %	RMSE (millis)	Under-predict %	RMSE (millis)
auth	17%	1755	30%	968
customer	32%	751	33%	665
flight	19%	978	15%	5604
booking	21%	1577	11%	5697

Prediction Error: Auth

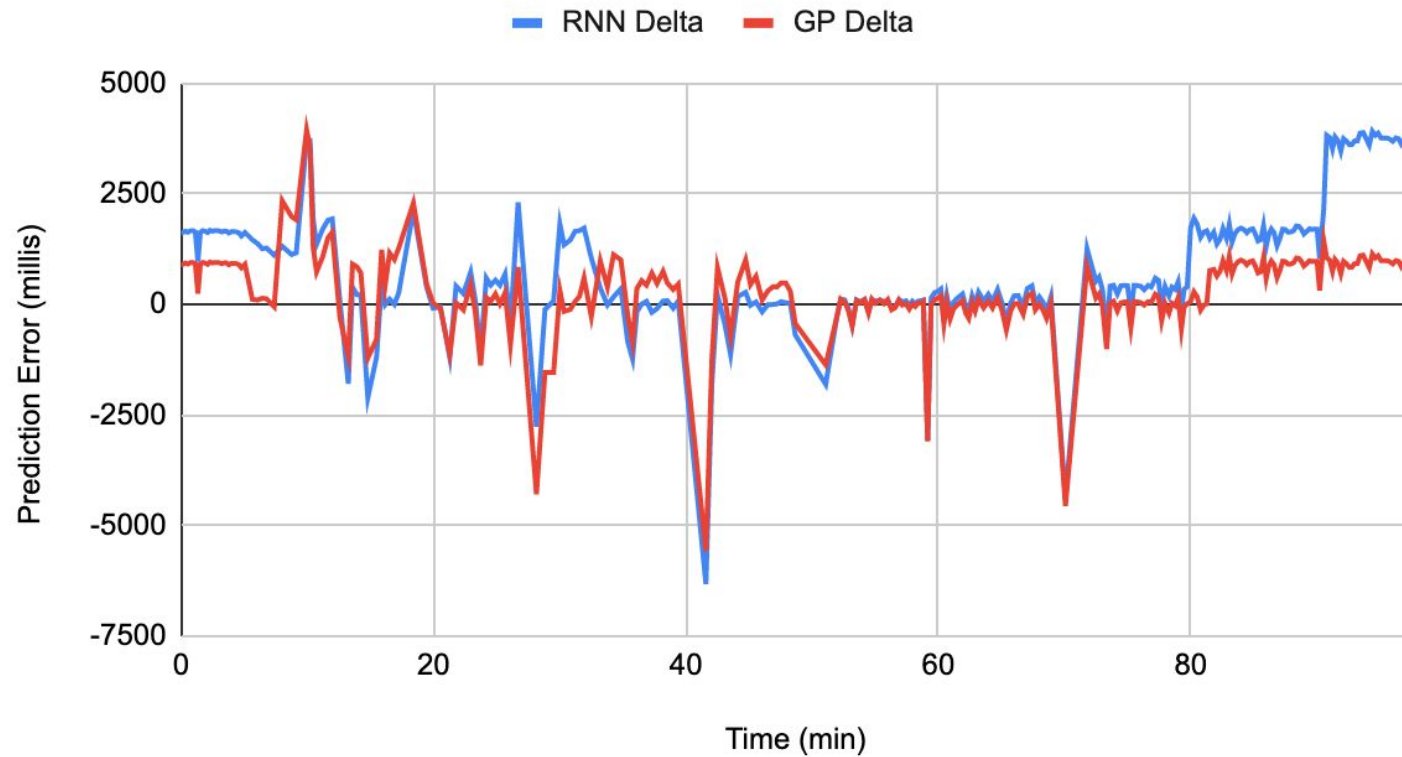


Table 1: Performance Metrics of Each Technique by Endpoint

Endpoint	RNN		GP	
	Under-predict %	RMSE (millis)	Under-predict %	RMSE (millis)
auth	17%	1755	30%	968
customer	32%	751	33%	665
flight	19%	978	15%	5604
booking	21%	1577	11%	5697

Prediction Error: Customer




Table 1: Performance Metrics of Each Technique by Endpoint

Endpoint	RNN		GP	
	Under-predict %	RMSE (millis)	Under-predict %	RMSE (millis)
auth	17%	1755	30%	968
customer	32%	751	33%	665
flight	19%	978	15%	5604
booking	21%	1577	11%	5697

Prediction Error: Flight




Table 1: Performance Metrics of Each Technique by Endpoint

Endpoint	RNN		GP	
	Under-predict %	RMSE (millis)	Under-predict %	RMSE (millis)
auth	17%	1755	30%	968
customer	32%	751	33%	665
flight	19%	978	15%	5604
booking	21%	1577	11%	5697


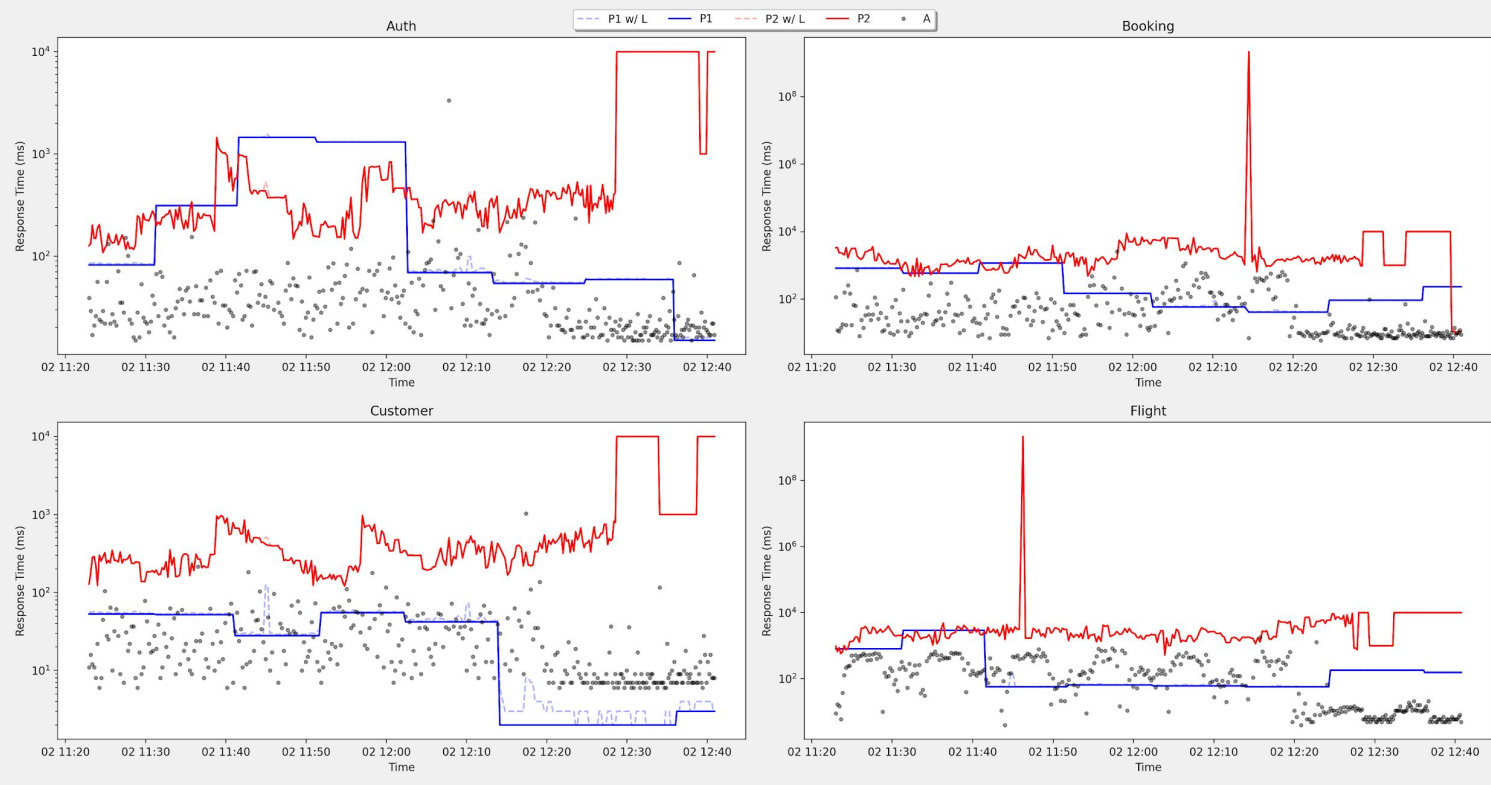
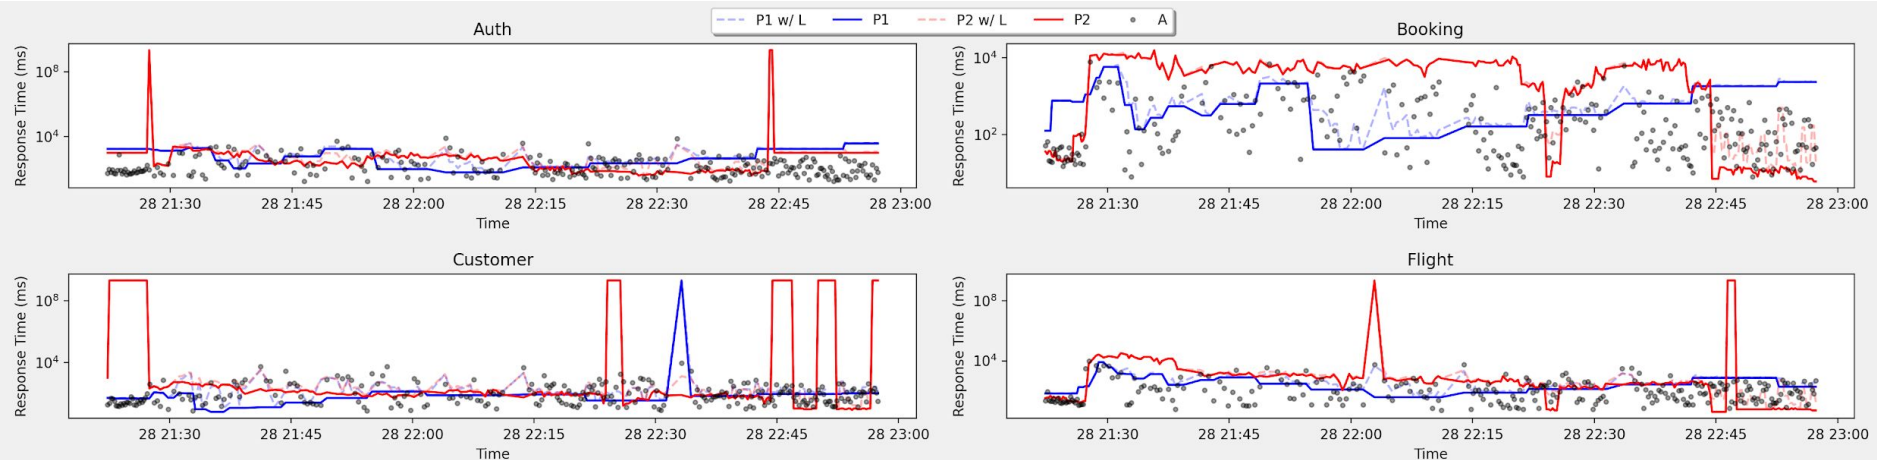


Table 1: Performance Metrics of Each Technique by Endpoint

Endpoint	RNN		GP	
	Under-predict %	RMSE (millis)	Under-predict %	RMSE (millis)
auth	17%	1755	30%	968
customer	32%	751	33%	665
flight	19%	978	15%	5604
booking	21%	1577	11%	5697



Future Direction

- Better prediction models
- Differentiating types of clients
- Granular endpoint
- Auto endpoint detection in broker
- Auto scaling resources
- Add more configurability
- Eliminate the additional network call?

